

Aza: AI-Powered Self-Service Resolution Hub

Amazon India

How a user review analysis of 13 reviews uncovered a ₹ 43 billion opportunity and led to a live deployed AI product built in under one week.

Role	Timeline	Status	Type
AI Product Manager	< 1 week	Live [Vercel]	Solo Project

01 — Executive Summary

Amazon already has a chatbot. Users are still complaining about it in 2025. Still being forced to convince a bot to connect them to a human. Still waiting 30-40 minutes for simple resolutions. This is not a case study about building something Amazon doesn't have. It's a study in what a better version looks like specifically around natural resolution, instant wallet refunds, and honest fraud thinking.

What Makes This Different from a Typical PM Case Study

Amazon already has a chatbot. They rebuilt it in 2024. User satisfaction improved by only 5%. Users as recently as 2025 still report having to convince the bot to connect them to a human for simple issues.

This is not a case study about building something Amazon doesn't have rather what a better version would look like.

- The live product is deployed and interactive. The business case uses stated assumptions, clearly labelled.
- The solution includes a fraud risk badge and proactive risk logic. The UI is live. The detection scoring has a bug, it stays 0/100. The logic and thinking behind it is documented here as a product decision, not a shipped feature.
- Every product decision is explicitly tied to a business metric or user outcome.

02 — The Problem

I began by doing what every PM does before building anything — reading real user reviews. I analyzed 10+ Amazon India user reviews and found the same four failure patterns repeating consistently.

Pain Point	Frequency	Business Risk
Long wait times— up to 1 hour on hold for simple queries	Very High	Purchase frequency drops, NPS collapse
Rude or unhelpful agents— unresolved after multiple calls	High	Trust erosion, churn, negative reviews
Forced to call back repeatedly for the same issue	High	Support cost multiplier, user frustration peak
Zero self service options— every issue needs human	Critical	Scales poorly, unsustainable at 150M users (number is an assumption as mentioned)

Root Cause — First principle breakdown

- **Surface problem** — Users call support for routine issues — waits, escalations, no resolution.
- **One level deeper** — Amazon has a chatbot but users still report it forces them to repeat themselves and escalate to humans for routine issues.
- **Two levels deeper** — The 2024 rebuild improved satisfaction by 5%, meaning the core problem of natural resolution and user trust was not addressed.
- **Root cause** — The existing chatbot optimises for deflection — keeping users away from human agents — not resolution. Users feel trapped, not helped. Aza is designed around the opposite goal — to resolve the issue completely, in the user’s own words without friction.

03 — Sizing the opportunity

Before designing any solution, I built a financial model to confirm this problem is worth solving at scale.

150M Active users	300M Annual contact	₹ 67.5B Annual support cost	₹ 43.2B Savings at 80% automation
-----------------------------	-------------------------------	---------------------------------------	---

Basis: 150M x 2contacts/year x ₹ 225/call (assumed)= ₹ 67.5B. At 80% automation rate (conservative for routine issues): ₹ 43.2B annual saving.

Conservative (50%) ₹ 21.6B \$ 260M	Base case (70%) ₹ 30.2B \$ 365M	Optimistic (80%) ₹ 43.2B \$ 520M
---	--	---

04 — Job-to-be-done Analysis

I mapped 4 primary user jobs identified directly from review patterns. Each scenario frames the user’s real goal, current friction preventing it, and the outcome they need.

01 Missing package	When: Order shows “delivered” but item never arrived Friction: Up to 1 hour on hold with delivery pattern — often unresolved Job: Instantly report issue and receive refund or replacement within 2 minutes Outcome: <i>Resolves in 2 minutes vs. wasting an entire afternoon</i>
02 Return an Item	When: Product doesn’t meet expectations or arrives damaged Friction: Return flow is confusing, requires calling support to initiate Job: Schedule a free return pickup automatically with zero phone calls Outcome: <i>No time wasted explaining problems to multiple different agents.</i>
03 Cancel an Order	When: User changes their mind before item has shipped Friction: Agent actively tries to dissuade cancellation — confrontational experience Job: Cancel instantly with one message, receive immediate confirmation Outcome: <i>Full autonomy over purchase decisions without agent friction</i>
04 Delivery change	When: User needs to update address or reschedule delivery window Friction: No option exists in the app — must call support to make any changes. Job: Update delivery details instantly through the app or chat Outcome: <i>Control over the delivery experience without human intervention</i>

05 — The Solution

Aza is an AI customer service assistant embedded directly into an Amazon India interface. It handles the top 80% of support contacts through natural, context-aware conversation — no hold times, no scripts, no human agents needed for routine issues.

Feature Architecture

Intent Detection Engine	Reads the full user message before responding. Extracts contexts, emotional tone, and issue type. Never asks users to repeat information they have already provided	Core
--------------------------------	---	-------------

Instant wallet refund	Process refund to Amazon pay wallet within 2 minutes of pickup confirmation. Money stays within Amazon's ecosystem — 73% of wallet credits are re-spent on amazon within 72 hours.	High Value
------------------------------	--	-------------------

Return & Replacement Flow	Schedules free pickup automatically. Offers refund or replacement with a single choice. Zero phone calls required end-to-end.	Core
--------------------------------------	---	-------------

Empathy Detection	Identifies frustrated or angry tone (including all-caps message). Always leads with genuine empathy before any action. Varies phrasing to avoid sounding scripted.	UX Quality
--------------------------	--	-------------------

Session Memory	Track conversation state throughout the session. Recognises when an issue is resolved and gives warm closing responses. Never asks about resolved issues again.	UX Quality
-----------------------	---	-------------------

Fraud Risk Scoring	A 0-100 risk badge is visible in the product. The product thinking behind it is real: session level scoring, red flags for refund-first behaviour, green flags for detailed complaints. Implementation is incomplete.	Trust & safety
---------------------------	---	---------------------------

Proactive Risk Prediction	The product logic is defined: damaged items should flag quality team, repeat issues should trigger goodwill credit. This is not implemented in the live product — it is a product decision documented here to show systems thinking beyond the core chatbot	Business Value
----------------------------------	---	-----------------------

Signature Product Decision: Instant Wallet Refund

The Business Logic Behind It

Standard refund: 3-5 business days back to the user’s bank account.
Aza’s approach: Instant credit to Amazon pay wallet, 2 minutes after the pickup confirmation.

Why this is better for Amazon — not just for users:

- The money never leaves the Amazon ecosystem.
- 73% of amazon pay credits are spent back on Amazon within 72 hours.
- A user who just had a bad experience immediately re-engages with the platform.
- Amazon converts a cost centre (refund processing) into a retention and revenue loop

This is the kind of decision that separates product thinking from feature thinking.

06 — PM Framework Applied

Job-to-be-done • Users don’t want to call support — they want resolution. • Mapped 4 JTBD scenarios directly from review data. • Every feature serves a specific, named user job. • Success = job completion rate, not conversation turns.	First Principles Thinking • What is the absolute minimum to resolve each issue? • Refund: order verification + payment trigger — fully automatable. • Cancellation: Ship status check + confirmation — fully automatable
---	---

Proactive Risk — Product thinking

- Damaged item: should flag quality team.
- Repeat issue: should trigger goodwill credit.
- Overdue order: should offer resolution proactively.
- Not implemented in live product: documented as product decision.
- Shows thinking beyond chatbot UX into ops and quality loops

Fraud Risk — Product Thinking

- The score should evaluate full session, not just individual messages.
- Red flag: First message is “refund” without context.
- Green flag: user describes real problem with detail
- Bug: Score currently stays 0 – logic does not trigger
- Next step: fix the use effect connection to messages state.

07 — Success Metrics Framework

North Star Metric — Time to Resolution

Definition: Average time from user first message to confirmed issue resolution.

Current baseline: 30-40 minutes (average phone call includes hold time).

Target with Aza: Under 2 minutes for 80% of issues.

Goal: 95% reduction in average resolution time within 6 months of launch.

Metric	Before Aza	After Aza	Impact
Average resolution time	30-40 minutes	< 2 minutes	95% faster
Cost per resolution	₹ 225/call	₹25/AI chat	89% reduction
AI containment rate	0% (no AI)	80% target	+80% points
CSAT Score	2.1/5 stars	4.5/5 stars	+114% improvement
First Contact Resolution	30%	95% target	+217% improvement
Annual cost saving	₹0 automated	₹43.2B	At 80% automation scale

08 — Key Product Learnings

1. The opportunity was not to build something new — It was to remove a bottleneck from something that already existed

Amazon had a return process. It has refunds, order tracking. The opportunity was not a missing feature — it was the missing connection between users and those features without humans in the middle.

2. Instant wallet refund are not UX feature — they are business strategy

Routing refunds to Amazon pay instead of a bank account costs Amazon nothing. The money stays on the platform. But it converts a frustrated, churning user back into an active buyer within hours. That is a retention mechanism disguised as a customer service improvement.

3. Fraud detection logic is harder to implement than to design

I designed clear scoring rules: session level evaluation, red flags for refund-first behaviour with no context, green flags for detailed genuine complaints. The UI badge is live. But the detection does not trigger — the score stays at 0. The bug is in how the scoring function connects to the message state. This taught me something real: Product thinking and code are two different things, and a PM needs to understand both well enough to catch the gap.

4. A chatbot that makes users repeat themselves fails the core job

The most critical failure of AI support tools is not understanding the context already provided. If a user describes their problem and the bot responds with ‘can you tell what happened?’ — the bot has failed. Intent detection must read the entire message, not keyword-match a fragment of it.

5. A live deployed product beats a perfect case study every time

I shipped a working, deployed product in under a week using AI tools — no prior coding background. A product a recruiter can actually open and interact with is worth more than 20 slides about a product that doesn’t exist.

09 — Honest Limitations & What I'd Do Next

Current Limitations

- The 10+ reviews sample is exploratory — needs larger validation.
- All order data is simulated.
- 80% automation rate is a hypothesis, not yet validated.
- The fraud score badge is live but stays at 0.
- Proactive risk logic defined but not implemented in live products.

Next Steps If Resourced

- Validate with 200+ reviews and 10+ real interviews.
- Integrate live order management and payment APIs.
- A/B Test: Instant wallet refund vs bank refund conversion.
- Build real-time PM metrics dashboard
- Add voice input for low-literacy user segment.

Live product: <https://akankshamohota30-aza-resolve-ai.vercel.app/>

Github: <https://github.com/akankshamohota30/resolve-ai>